

七月在线电商推荐系统项目特训营结业考试参考答案

By 魏天保

一、理论题（40分）

1、协同过滤算法是指基于用户行为数据设计的推荐算法，常见的有UserCF和ItemCF两种，请你分别用一句话概括这两种算法的核心思想。（10分）

答案：UserCF的核心思想是：给用户推荐和他兴趣相似的其他用户喜欢的物品。ItemCF的核心思想是：给用户推荐和其过去感兴趣的物品相似的物品。

2、矩阵分解属于协同过滤算法吗？（5分）

A 属于

B 不属于

答案：A

3、请列举出基于内容的推荐的两个优点（6分）

答：（1）没有物品冷启动的问题（2）推荐结果直观，容易理解

4、请说出推荐系统中评分预测和TopN推荐的评价指标，各列举两种。（4分）

答案：评分预测：RMSE、MAE；TopN推荐：Precision、Recall。

5、写出RMSE的计算公式，并说明公式中每个变量的含义。（5分）

答案：

$$RMSE = \sqrt{\frac{1}{m} \sum_{i=1}^m (y_i - \hat{y}_i)^2}$$

m是测试集样本个数， y_i 是第i个测试样本真实值， \hat{y}_i 是第i个测试样本预测值

6、Deep AutoEncoder 的核心思想是什么。（10分）

答案：核心的思想是通过反向传播训练网络，最小化输入和输出的误差，从而完成对未知的电影的预估。

二、代码题（60分）

7、请按要求完成基于 tensorflow 的 Deep AutoEncoder 算法。（共30分）

In [1]:

```
1 import warnings
2 warnings.filterwarnings('ignore')
3 import numpy as np
4 import pandas as pd
5 import tensorflow as tf
```

加载数据

In [2]:

```
1 df = pd.read_csv('ratings.dat', sep='\t', names=['user', 'item', 'rating', 'timestamp'], header=0)
```

(1) 去掉评分数据 df 中的 “timestamp” 这一列，并输出 df 前五行（5分）

In [3]:

```
1 df = df.drop('timestamp', axis=1)
2 df.head()
```

Out[3]:

	user	item	rating
0	1	1193	5
1	1	661	3
2	1	914	3
3	1	3408	4
4	1	2355	5

查看用户数和电影数

In [4]:

```
1 num_items = df.item.nunique()
2 num_users = df.user.nunique()
3 print("USERS: {} ITEMS: {}".format(num_users, num_items))
```

USERS: 6040 ITEMS: 3706

(2) 对 df 的 “rating” 列做 Normalization，展示 df 前 5 行（5分）

In [5]:

```

1 from sklearn import preprocessing
2 r = df['rating'].values.astype(float)
3 min_max_scaler = preprocessing.MinMaxScaler()
4 x_scaled = min_max_scaler.fit_transform(r.reshape(-1,1))
5 df_normalized = pd.DataFrame(x_scaled)
6 df['rating'] = df_normalized
7 df.head()

```

Out[5]:

	user	item	rating
0	1	1193	1.00
1	1	661	0.50
2	1	914	0.50
3	1	3408	0.75
4	1	2355	1.00

(3) 把 df 转成 user-item 矩阵，存储在变量 matrix 中，用 0 填充矩阵的缺失值，从 matrix 中取出 user5 对 item6 的评分值（5分）

In [6]:

```

1 matrix = df.pivot(index='user', columns='item', values='rating')
2 matrix.fillna(0, inplace=True)
3 matrix.loc[5,6]

```

Out[6]:

0.25

获取用户和物品列表

In [7]:

```

1 users = matrix.index.tolist()
2 items = matrix.columns.tolist()

```

将 matrix 从 dataframe 转换成 numpy 数组

In [8]:

```
1 matrix = matrix.as_matrix()
```

网络超参数

In [9]:

```
1 num_input = num_items
2 num_hidden_1 = 10
3 num_hidden_2 = 5
```

隐层的变量初始化

In [10]:

```
1 weights = {
2     'encoder_h1': tf.Variable(tf.random_normal([num_input, num_hidden_1], dtype=tf.float64)),
3     'encoder_h2': tf.Variable(tf.random_normal([num_hidden_1, num_hidden_2], dtype=tf.float64)),
4     'decoder_h1': tf.Variable(tf.random_normal([num_hidden_2, num_hidden_1], dtype=tf.float64)),
5     'decoder_h2': tf.Variable(tf.random_normal([num_hidden_1, num_input], dtype=tf.float64)),
6 }
7
8 biases = {
9     'encoder_b1': tf.Variable(tf.random_normal([num_hidden_1], dtype=tf.float64)),
10    'encoder_b2': tf.Variable(tf.random_normal([num_hidden_2], dtype=tf.float64)),
11    'decoder_b1': tf.Variable(tf.random_normal([num_hidden_1], dtype=tf.float64)),
12    'decoder_b2': tf.Variable(tf.random_normal([num_input], dtype=tf.float64)),
13 }
```

(4) 构建 encoder 和 decoder (5分)

In [11]:

```
1 def encoder(x):
2     layer_1 = tf.nn.relu(tf.add(tf.matmul(x, weights['encoder_h1']), biases['encoder_b1']))
3     layer_2 = tf.nn.relu(tf.add(tf.matmul(layer_1, weights['encoder_h2']), biases['encoder_b2']))
4     return layer_2
5
6 def decoder(x):
7     layer_1 = tf.nn.relu(tf.add(tf.matmul(x, weights['decoder_h1']), biases['decoder_b1']))
8     layer_2 = tf.nn.relu(tf.add(tf.matmul(layer_1, weights['decoder_h2']), biases['decoder_b2']))
9     return layer_2
```

构建整个模型

In [12]:

```
1 X = tf.placeholder(tf.float64, [None, num_input])
2 encoder_op = encoder(X)
3 decoder_op = decoder(encoder_op)
```

预测 y 值和真实 y 值

In [13]:

```
1 y_pred = decoder_op
2 y_true = X
```

定义损失函数和优化器

In [14]:

```
1 loss = tf.losses.mean_squared_error(y_true, y_pred)
2 optimizer = tf.train.RMSPropOptimizer(0.03).minimize(loss)
```

定义评估准则

In [15]:

```
1 eval_x = tf.placeholder(tf.int32, )
2 eval_y = tf.placeholder(tf.int32, )
3 pre, pre_op = tf.metrics.precision(labels=eval_x, predictions=eval_y)
```

变量初始化

In [16]:

```
1 init = tf.global_variables_initializer()
2 local_init = tf.local_variables_initializer()
```

在 session 中 run

In [17]:

```

1 predictions = pd.DataFrame()
2 with tf.Session() as session:
3     epochs = 100
4     batch_size = 250
5
6     session.run(init)
7     session.run(local_init)
8
9     num_batches = int(matrix.shape[0] / batch_size)
10    matrix = np.array_split(matrix, num_batches)
11
12    for i in range(epochs):
13
14        avg_cost = 0
15
16        for batch in matrix:
17            _, l = session.run([optimizer, loss], feed_dict={X: batch})
18            avg_cost += l
19
20        avg_cost /= num_batches
21
22        print("Epoch: {} Loss: {}".format(i + 1, avg_cost))
23
24    print("Predictions...")
25
26    matrix = np.concatenate(matrix, axis=0)
27
28    preds = session.run(decoder_op, feed_dict={X: matrix})
29
30    predictions = predictions.append(pd.DataFrame(preds))
31
32    predictions = predictions.stack().reset_index(name='rating')
33    predictions.columns = ['user', 'item', 'rating']
34    predictions['user'] = predictions['user'].map(lambda value: users[value])
35    predictions['item'] = predictions['item'].map(lambda value: items[value])

```

```

Epoch: 1 Loss: 31.39658373594284
Epoch: 2 Loss: 1.4016507441798847
Epoch: 3 Loss: 0.6289405835171541
Epoch: 4 Loss: 0.47894836962223053
Epoch: 5 Loss: 0.41464972496032715
Epoch: 6 Loss: 0.2824529707431793
Epoch: 7 Loss: 0.12261824092517297
Epoch: 8 Loss: 0.04141972423531115
Epoch: 9 Loss: 0.023202844196930528
Epoch: 10 Loss: 0.020845729508437216
Epoch: 11 Loss: 0.020702326049407322
Epoch: 12 Loss: 0.02070566701392333
Epoch: 13 Loss: 0.020706364419311285
Epoch: 14 Loss: 0.020706464264852304
Epoch: 15 Loss: 0.02070651645772159
Epoch: 16 Loss: 0.020706516418916483
Epoch: 17 Loss: 0.020706516263696056
Epoch: 18 Loss: 0.020706516263696056
Epoch: 19 Loss: 0.02070651634130627
Epoch: 20 Loss: 0.02070651634130627
Epoch: 21 Loss: 0.020706516263696056
Epoch: 22 Loss: 0.02070651634130627

```

Epoch: 23 Loss: 0.02070651634130627
Epoch: 24 Loss: 0.02070651634130627
Epoch: 25 Loss: 0.02070651634130627
Epoch: 26 Loss: 0.02070651634130627
Epoch: 27 Loss: 0.02070651634130627
Epoch: 28 Loss: 0.020706516418916483
Epoch: 29 Loss: 0.02070651634130627
Epoch: 30 Loss: 0.020706516263696056
Epoch: 31 Loss: 0.020706516263696056
Epoch: 32 Loss: 0.020706516418916483
Epoch: 33 Loss: 0.02070651634130627
Epoch: 34 Loss: 0.02070651634130627
Epoch: 35 Loss: 0.02070651634130627
Epoch: 36 Loss: 0.020706516263696056
Epoch: 37 Loss: 0.02070651634130627
Epoch: 38 Loss: 0.02070651634130627
Epoch: 39 Loss: 0.02070651634130627
Epoch: 40 Loss: 0.02070651634130627
Epoch: 41 Loss: 0.02070651634130627
Epoch: 42 Loss: 0.02070651634130627
Epoch: 43 Loss: 0.02070651634130627
Epoch: 44 Loss: 0.02070651634130627
Epoch: 45 Loss: 0.020706516263696056
Epoch: 46 Loss: 0.02070651634130627
Epoch: 47 Loss: 0.02070651634130627
Epoch: 48 Loss: 0.020706516263696056
Epoch: 49 Loss: 0.020706516263696056
Epoch: 50 Loss: 0.02070651634130627
Epoch: 51 Loss: 0.02070651634130627
Epoch: 52 Loss: 0.02070651634130627
Epoch: 53 Loss: 0.020706516263696056
Epoch: 54 Loss: 0.02070651634130627
Epoch: 55 Loss: 0.020706516263696056
Epoch: 56 Loss: 0.02070651634130627
Epoch: 57 Loss: 0.020706516263696056
Epoch: 58 Loss: 0.020706516263696056
Epoch: 59 Loss: 0.02070651634130627
Epoch: 60 Loss: 0.020706516263696056
Epoch: 61 Loss: 0.02070651634130627
Epoch: 62 Loss: 0.02070651634130627
Epoch: 63 Loss: 0.020706516263696056
Epoch: 64 Loss: 0.020706516263696056
Epoch: 65 Loss: 0.02070651634130627
Epoch: 66 Loss: 0.020706516263696056
Epoch: 67 Loss: 0.020706516263696056
Epoch: 68 Loss: 0.020706516418916483
Epoch: 69 Loss: 0.020706516263696056
Epoch: 70 Loss: 0.020706516263696056
Epoch: 71 Loss: 0.02070651634130627
Epoch: 72 Loss: 0.02070651634130627
Epoch: 73 Loss: 0.020706516418916483
Epoch: 74 Loss: 0.020706516418916483
Epoch: 75 Loss: 0.02070651634130627
Epoch: 76 Loss: 0.02070651634130627
Epoch: 77 Loss: 0.020706516263696056
Epoch: 78 Loss: 0.02070651634130627
Epoch: 79 Loss: 0.020706516263696056
Epoch: 80 Loss: 0.02070651634130627
Epoch: 81 Loss: 0.02070651634130627
Epoch: 82 Loss: 0.02070651634130627
Epoch: 83 Loss: 0.02070651634130627

```
Epoch: 84 Loss: 0.02070651634130627
Epoch: 85 Loss: 0.020706516418916483
Epoch: 86 Loss: 0.02070651634130627
Epoch: 87 Loss: 0.020706516263696056
Epoch: 88 Loss: 0.020706516263696056
Epoch: 89 Loss: 0.020706516418916483
Epoch: 90 Loss: 0.02070651634130627
Epoch: 91 Loss: 0.02070651634130627
Epoch: 92 Loss: 0.020706516263696056
Epoch: 93 Loss: 0.020706516263696056
Epoch: 94 Loss: 0.020706516263696056
Epoch: 95 Loss: 0.020706516418916483
Epoch: 96 Loss: 0.02070651634130627
Epoch: 97 Loss: 0.02070651634130627
Epoch: 98 Loss: 0.02070651634130627
Epoch: 99 Loss: 0.020706516263696056
Epoch: 100 Loss: 0.020706516263696056
Predictions...
```

(5) 为用户42计算top10的推荐结果 (10分)

- top10 的筛选依据是评分预测最高的 top10
- 要求推荐的 top10 是该用户没有看过的电影
- 最后输出推荐的 top10 电影的 itemid

In [18]:

```
1 seen_movies = df[df.user==42].item.values
2 rec_for_42 = predictions[(predictions.user==42)&(~predictions.item.isin(seen_movies))]
3 top10 = rec_for_42.sort_values('rating',ascending=False).head(10).item.values
4 print(top10)
```

```
[2396 110    1  912   50 1247 1259  111 1225 2324]
```

8、请你实现基于Keras的协同深度学习算法。 (共30分)

In [1]:

```
1 import numpy as np
2 from pandas import read_csv
3 from sklearn.preprocessing import LabelEncoder
4 from keras.layers import Input, Embedding, Dot, Flatten, Dense, Dropout, Lambda, Add
5 from keras.layers.noise import GaussianNoise
6 from keras.initializers import RandomUniform, RandomNormal
7 from keras.models import Model
8 from keras.regularizers import l2
9 from keras import optimizers
10 from keras import backend as K
11 from keras.engine.topology import Layer
```

Using TensorFlow backend.

(1) 请你完成 movie_map 这个函数 (5分)

- 功能：使用自然数对 movies.dat 中的 movieid 进行重编码
- 返回值：movieid 到编码 id 的字典
- 测试：计算编码字典，存入变量 d (后面有用)，输出 d[520]

In [2]:

```
1 def movie_map(file='movies.dat'):
2     movies = read_csv(file, sep='::', header=None, engine='python')
3     value = movies[0].unique()
4     index = range(movies[0].nunique())
5     return dict(zip(value, index))
6
7 d = movie_map()
8 d[520]
```

Out[2]:

516

(2) 请你完成 read_ratings 这个函数 (5分)

- 要求返回数据类型为 float32 的二维 numpy 数组：每一行表示用户评分记录 [user,item,rating]
- 原始数据中的 itemid 需要使用编码字典 d 进行重编码
- rating 值需要进行幅度缩放，缩放区间为[0,1]
- 测试：调用 read_ratings ,将返回值存入变量 ratings （后面有用） ， 并输出 ratings[:3]

In [3]:

```
1 def read_ratings(file='ratings.dat'):
2     rating_mat = list()
3     with open(file) as fp:
4         for line in fp:
5             line = line.strip().split('\t')
6             user, item, rating = line[0], d[int(line[1])], int(line[2])/5
7             rating_mat.append( [user, item, rating] )
8     return np.array(rating_mat).astype('float32')
9
10 ratings = read_ratings()
11 ratings[:3]
```

Out[3]:

```
array([[1.000e+00, 1.176e+03, 1.000e+00],
       [1.000e+00, 6.550e+02, 6.000e-01],
       [1.000e+00, 9.020e+02, 6.000e-01]], dtype=float32)
```

(3) 请你完成 train_test_split 这个函数 (5分)

- 功能：将评分数据 ratings 划分为训练集和测试集两部分
- 输入：函数 read_ratings 的返回值 ratings
- 输出：训练集 train_mat，测试集 test_mat
- 要求：训练集占 80%，测试集占 20%
- 测试：调用该函数，返回值存入 train_mat, test_mat （后面有用） ， 并展示训练集和测试集的样本数量

In [4]:

```

1 def train_test_split(ratings):
2
3     data_num = len(ratings)
4     train_mat = ratings[:int(data_num*0.8)]
5     test_mat = ratings[int(data_num*0.8):]
6
7     return train_mat, test_mat
8
9 train_mat, test_mat = train_test_split(ratings)
10 train_mat.shape[0], test_mat.shape[0]

```

Out[4]:

(800167, 200042)

(4) 请你完成 read_item 这个函数 (5分)

- 功能：将 movies.dat 的三列特征进行重编码
- 要求：
 - 第 1 列特征使用编码字典 d 进行编码
 - 第 2、3 列特征使用 LabelEncoder() 编码
 - 返回值有两个，第一个是二维 numpy 数组，也就是一个矩阵：每一行对应原始数据 movies.dat 的一行记录；第二个是特征维度，也就是矩阵的列数
 - 调用该函数，返回值存入 item_mat,item_feat_dim （后面有用），打印 item_mat[:3] 和 item_feat_dim

In [5]:

```

1 def read_item(file='movies.dat'):
2
3     item = read_csv(file, sep='::', header=None, engine='python')
4     item[0] = item[0].apply(lambda x:d[int(x)])
5     item[1] = LabelEncoder().fit_transform(item[1])
6     item[2] = LabelEncoder().fit_transform(item[2])
7     item_mat = item.as_matrix()
8     item_feat_dim = item_mat.shape[1]
9
10    return item_mat,item_feat_dim
11
12 item_mat,item_feat_dim = read_item()
13 print(item_mat[:3])
14 print(item_feat_dim)

```

```

[[ 0 3574 145]
 [ 1 1858 115]
 [ 2 1483 207]]

```

3

CollaborativeDeepLearning 实现

In [6]:

```

1  class CollaborativeDeepLearning:
2
3      def __init__(self, item_mat, hidden_layers):
4          """
5          hidden_layers = a list of three integer indicating the embedding dimension of autoencoder
6          item_mat = item feature matrix with shape (# of item, # of item features)
7          """
8          assert(len(hidden_layers)==3)
9          self.item_mat = item_mat
10         self.hidden_layers = hidden_layers
11         self.item_dim = hidden_layers[0]
12         self.embedding_dim = hidden_layers[-1]
13
14     def pretrain(self, lamda_w=0.1, encoder_noise=0.1, dropout_rate=0.1, activation='sigmoid',
15                 """
16                 layer-wise pretraining on item features (item_mat)
17                 """
18                 self.trained_encoders = []
19                 self.trained_decoders = []
20                 X_train = self.item_mat
21
22                 for input_dim, hidden_dim in zip(self.hidden_layers[:-1], self.hidden_layers[1:]):
23
24                     pretrain_input = Input(shape=(input_dim,))
25                     encoded = GaussianNoise(stddev=encoder_noise)(pretrain_input)
26                     encoded = Dropout(dropout_rate)(encoded)
27                     encoder = Dense(hidden_dim, activation=activation, kernel_regularizer=l2(lamda_w),
28                                     decoder = Dense(input_dim, activation=activation, kernel_regularizer=l2(lamda_w),
29
30                     # autoencoder
31                     ae = Model(inputs=pretrain_input, outputs=decoder)
32
33                     # encoder
34                     ae_encoder = Model(inputs=pretrain_input, outputs=encoder)
35
36                     # decoder
37                     encoded_input = Input(shape=(hidden_dim,))
38
39                     decoder_layer = ae.layers[-1] # the last layer
40                     ae_decoder = Model(encoded_input, decoder_layer(encoded_input))
41
42                     ae.compile(loss='mse', optimizer='rmsprop')
43                     ae.fit(X_train, X_train, batch_size=batch_size, epochs=epochs, verbose=2)
44
45                     self.trained_encoders.append(ae_encoder)
46                     self.trained_decoders.append(ae_decoder)
47                     X_train = ae_encoder.predict(X_train)
48
49     def finetune(self, train_mat, test_mat, lamda_u=0.1, lamda_v=0.1, lamda_n=0.1, lr=0.001, b
50     """
51     Fine-tuning with rating prediction
52     """
53     num_user = int( max(train_mat[:,0].max(), test_mat[:,0].max()) + 1 )
54     num_item = int( max(train_mat[:,1].max(), test_mat[:,1].max()) + 1 )
55
56     # item autoencoder
57     itemfeat_InputLayer = Input(shape=(self.item_dim,), name='item_feat_input')
58     encoded = self.trained_encoders[0](itemfeat_InputLayer)
59     encoded = self.trained_encoders[1](encoded)

```

```

60     decoded = self.trained_decoders[1](encoded)
61     decoded = self.trained_decoders[0](decoded)
62
63     # user embedding
64     user_InputLayer = Input(shape=(1,), dtype='int32', name='user_input')
65     user_EmbeddingLayer = Embedding(input_dim=num_user, output_dim=self.embedding_dim,
66                                     input_length=1, name='user_embedding',
67                                     embeddings_regularizer=l2(lamda_u),
68                                     embeddings_initializer=RandomNormal(mean=0, stddev=1))
69     user_EmbeddingLayer = Flatten(name='user_flatten')(user_EmbeddingLayer)
70
71     # item embedding
72     item_InputLayer = Input(shape=(1,), dtype='int32', name='item_input')
73     item_OffsetVector = Embedding(input_dim=num_item, output_dim=self.embedding_dim,
74                                   input_length=1, name='item_offset_vector',
75                                   embeddings_regularizer=l2(lamda_v),
76                                   embeddings_initializer=RandomNormal(mean=0, stddev=1))(i
77     item_OffsetVector = Flatten(name='item_flatten')(item_OffsetVector)
78     item_EmbeddingLayer = Add()([encoded, item_OffsetVector])
79
80     # rating prediction
81     dotLayer = Dot(axes = -1, name='dot_layer')([user_EmbeddingLayer, item_EmbeddingLayer]
82
83     self.cdl_model = Model(inputs=[user_InputLayer, item_InputLayer, itemfeat_InputLayer],
84                             self.cdl_model.compile(optimizer='rmsprop', loss=['mse', 'mse'], loss_weights=[1, lamda
85
86     train_user, train_item, train_item_feat, train_label = self.matrix2input(train_mat)
87     test_user, test_item, test_item_feat, test_label = self.matrix2input(test_mat)
88
89     model_history = self.cdl_model.fit([train_user, train_item, train_item_feat],
90                                       [train_label, train_item_feat],
91                                       epochs=epochs, batch_size=batch_size,
92                                       validation_data=([test_user, test_item, test_item_fe
93     return model_history
94
95     def matrix2input(self, rating_mat):
96         train_user = rating_mat[:, 0].reshape(-1, 1).astype(int)
97         train_item = rating_mat[:, 1].reshape(-1, 1).astype(int)
98         train_label = rating_mat[:, 2].reshape(-1, 1)
99         train_item_feat = [self.item_mat[train_item[x]][0] for x in range(train_item.shape[0])]
100        return train_user, train_item, np.array(train_item_feat), train_label
101
102    def build(self):
103        # rating prediction
104        prediction_layer = Dot(axes = -1, name='prediction_layer')([user_EmbeddingLayer, encod
105        self.model = Model(inputs=[user_InputLayer, itemfeat_InputLayer], outputs=[prediction_
106
107    def getRMSE(self, test_mat):
108        test_user, test_item, test_item_feat, test_label = self.matrix2input(test_mat)
109        pred_out = self.cdl_model.predict([test_user, test_item, test_item_feat])
110        return np.sqrt(np.mean(np.square(test_label.flatten() - pred_out[0].flatten())))

```

(5) 成功运行下面的代码，输出RMSE (10分)

In [7]:

```

1 model = CollaborativeDeepLearning(item_mat, [item_feat_dim, 16, 8])
2 model.pretrain(lamda_w=0.001, encoder_noise=0.3, epochs=10)
3 model_history = model.finature(train_mat, test_mat, lamda_u=0.01, lamda_v=0.1, lamda_n=0.1, lr=0.001)
4 model.getRMSE(test_mat)

```

Epoch 1/10

0s - loss: 3362782.6749

Epoch 2/10

0s - loss: 3362457.1563

Epoch 3/10

0s - loss: 3362184.4979

Epoch 4/10

0s - loss: 3361981.1230

Epoch 5/10

0s - loss: 3361836.0971

Epoch 6/10

0s - loss: 3361736.3306

Epoch 7/10

0s - loss: 3361669.6875

Epoch 8/10

0s - loss: 3361630.6741

Epoch 9/10

0s - loss: 3361608.1432

Epoch 10/10

0s - loss: 3361595.3847

Epoch 1/10

0s - loss: 0.2609

Epoch 2/10

0s - loss: 0.2068

Epoch 3/10

0s - loss: 0.1644

Epoch 4/10

0s - loss: 0.1371

Epoch 5/10

0s - loss: 0.1213

Epoch 6/10

0s - loss: 0.1124

Epoch 7/10

0s - loss: 0.1070

Epoch 8/10

0s - loss: 0.1039

Epoch 9/10

0s - loss: 0.1019

Epoch 10/10

0s - loss: 0.1005

Train on 800167 samples, validate on 200042 samples

Epoch 1/4

```

800167/800167 [=====] - 47s - loss: 320012.3428 - dot_layer
_loss: 0.7405 - model_3_loss: 3198583.5023 - val_loss: 311579.5984 - val_dot_layer_l
oss: 0.5741 - val_model_3_loss: 3115789.5686

```

Epoch 2/4

```

800167/800167 [=====] - 44s - loss: 319858.8802 - dot_layer
_loss: 0.4799 - model_3_loss: 3198583.4552 - val_loss: 311579.5988 - val_dot_layer_l
oss: 0.5745 - val_model_3_loss: 3115789.5686

```

Epoch 3/4

```

800167/800167 [=====] - 48s - loss: 319858.8803 - dot_layer
_loss: 0.4800 - model_3_loss: 3198583.4564 - val_loss: 311579.5697 - val_dot_layer_l
oss: 0.5741 - val_model_3_loss: 3115789.5686

```

Epoch 4/4

```
800167/800167 [=====] - 52s - loss: 319858.8798 - dot_layer  
_loss: 0.4800 - model_3_loss: 3198583.4517 - val_loss: 311579.5988 - val_dot_layer_l  
oss: 0.5745 - val_model_3_loss: 3115789.5686
```

Out[7]:

0.75795275